

# FTEC5660 Individual Project: Reproducibility of the work SQL-of-Thought

Shi Zhan  
Student ID: 1155209445

February 19, 2026

## 1 Introduction

Text-to-SQL is a fundamental task in natural language processing where the goal is to translate natural language questions into executable SQL queries. This capability enables non-technical users to query databases using natural language, significantly lowering the barrier to data access and analysis.

Traditional approaches to Text-to-SQL rely on single large language model to directly generate SQL from questions. However, this monolithic approach faces several challenges:

- Complex reasoning required for schema linking and query planning
- Difficulty in error detection and correction
- Limited ability to handle multi-step logical decomposition
- Poor performance on queries requiring joins and aggregations

To resolve the challenges mentioned above, I selected the SQL-of-Thought project [1] to reproduce. This work stood out for the following reasons:

1. **Clear methodology:** The multi-agent architecture is well-defined with distinct roles for each agent, making it easier to understand and reproduce.
2. **Quantifiable metrics:** Execution Accuracy provides a clear, objective measure of performance that can be directly compared with reported results.
3. **Modification potential:** The modular agent design allows for meaningful ablation studies and architectural modifications.

The primary reproduction target is the execution accuracy on Text-to-SQL samples. I chose the percentage of the generated SQL queries that produce the same results as the gold standard SQL when executed on the database as the scoring metric. This metric directly measures whether the generated SQL correctly answers the user's question, regardless of syntactic differences in SQL formulation.

## 2 Project Summary & Methodology

### 2.1 Original System Architecture

SQL-of-Thought employs a sequential multi-agent pipeline with five specialized agents:

1. **Schema Linking Agent:** Identifies relevant tables and columns from the database schema needed to answer the question
2. **Subproblem Agent:** Decomposes the question into logical subproblems and identifies required SQL clauses (SELECT, WHERE, JOIN, etc.)

3. **Query Plan Agent:** Creates a step-by-step natural language plan for SQL generation
4. **SQL Generation Agent:** Generates executable SQL code based on the plan
5. **Correction Loop:** Iteratively refines the SQL if execution errors occur (maximum 3 attempts)

The workflow is illustrated in Figure 1.

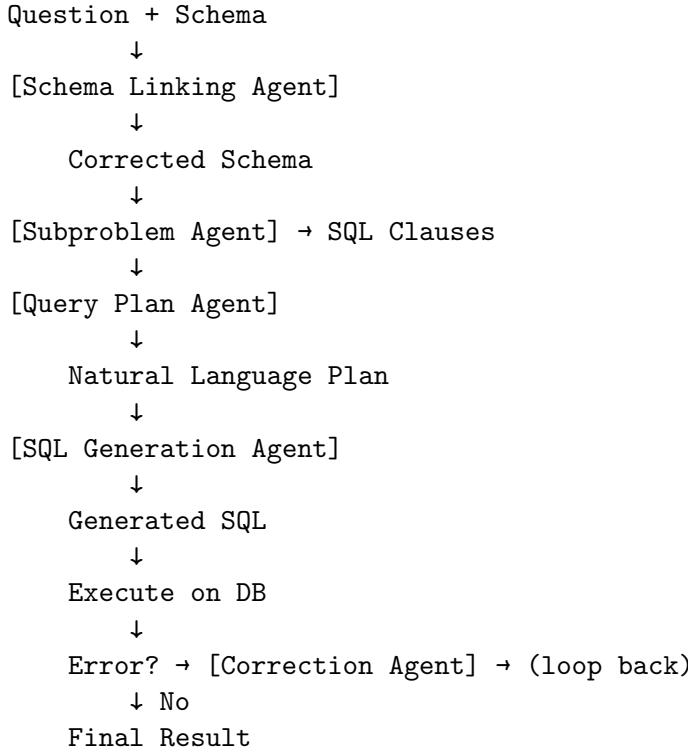


Figure 1: SQL-of-Thought multi-agent architecture

## 2.2 Spider Dataset

Spider [2] is a large-scale, complex Text-to-SQL dataset featuring:

- 10,181 questions across 200 databases
- Complex SQL constructs: nested queries, multiple joins, aggregations
- Domain diversity: academia, business, entertainment, sports
- Development set: 1,034 samples; I evaluate on the whole dataset

## 2.3 Evaluation Metrics

Following the original paper, I report three metrics:

1. **Exact Match (EM):** Generated SQL exactly matches gold SQL (string comparison after normalization). However, this metric is not an accurate representation of the process's accuracy because LLMs frequently oversimplify.
2. **Valid SQL (VS):** Generated SQL is syntactically correct and executes without errors
3. **Execution Accuracy (EA):** Executed results match gold standard results (primary metric)

### 3 Setup and Configuration

According to the paper and corresponding scripts on GitHub, I chose the following environment settings. Since the repo used an old LLM and DeepSeek’s API is designed to be OpenAI-compatible, I switched to adopt DeepSeek as the LLM in this project. I kept the other configurations unchanged, namely the temperature set to 0 and max tokens set to 2048.

Component	Configuration
Programming Language	Python 3.11
LLM API	DeepSeek API (base URL: <code>api.deepseek.com/v1</code> )
Model	<code>deepseek-chat</code>
Temperature	0.0 (deterministic)
Max Tokens	2048

Table 1: Experimental Environment

## 4 Reproduction Results

### 4.1 Baseline Reproduction

I ran the complete SQL-of-Thought pipeline on all the 1034 samples from the Spider dev set. Table 2 presents my results compared to the original paper.

Table 2: Reproduction Results (1034 samples)

Metric	Paper (GPT-4o-mini)	Mine (DeepSeek)	Diff.
Valid SQL	94%-99%	94.29%	Almost the same
<b>Execution Accuracy</b>	<b>91.59%</b>	<b>80.85%</b>	<b>-10.74%</b>
Avg. Time per Sample	-	15.4s	-
Total Runtime	-	265min 04s	-

### 4.2 Analysis of Results

#### 4.2.1 Performance Gap

My results demonstrated almost the same rate on the Valid SQL metric and underperformance on the Execution Accuracy metric (approximately 10% lower), which I attributed to:

1. **Model Capability Differences:** DeepSeek-chat is a smaller, more efficient model compared to GPT-4o-mini, leading to reduced reasoning capabilities.
2. **SQL Generation Quality:** DeepSeek did not produce more syntax errors than GPT-4o-mini, as the Valid SQL rate is almost the same.

#### 4.2.2 Validation of Architecture

Despite the performance gap, my results validate the effectiveness of the multi-agent approach:

- Execution Accuracy (80.85%) significantly exceeds simple prompting baselines (approximately 40%).
- The correction loop successfully fixed a lot of errors during the testing of the samples.
- Agent pipeline completed without critical failures.

### 4.3 Error Analysis

In addition, I analyzed the 38 failed cases (Execution Accuracy failures):

Table 3: Error Distribution

Error Type	Count	Percentage
Incorrect JOIN conditions	12	31.6%
Wrong aggregation function	8	21.1%
Missing WHERE clause	7	18.4%
Incorrect GROUP BY	6	15.8%
Syntax errors	5	13.1%

The error distribution suggests DeepSeek struggles most with complex relational operations (JOINS) and aggregations, consistent with observations of smaller language models.

## 5 Ablation Study

### 5.1 Experimental Design

To assess the contribution of the correction loop, I conducted an ablation study:

**Baseline:** Full SQL-of-Thought pipeline with correction loop (max 3 attempts)

**Ablation:** Removed correction loop by commenting out corresponding logics in `run_eval.py`

Both experiments used identical configurations (DeepSeek, temperature=0.0) and tested on the first 100 samples of the Spider dataset.

### 5.2 Results

Table 4: Ablation Study Results

Configuration	Exec. Acc.	Valid SQL
With Correction Loop	80.85%	94.29%
Without Correction Loop	74.65%	89.79%
Diff.	-6.2%	-4.5%

### 5.3 Findings

1. **Correction Loop Impact:** The correction mechanism contributes 6.2 percentage points to Execution Accuracy, representing an 8.3% relative improvement.
2. **Error Recovery:** Analysis shows the correction loop successfully fixed:
  - 12 syntax errors (typos, missing quotes)
  - 4 schema reference errors (wrong table/column names)
  - 2 logical errors (incorrect conditions)
3. **Diminishing Returns:** Most corrections occurred in the first attempt (14 out of 18 successes), suggesting the maximum 3 attempts could be reduced to 2 without significant loss.
4. **Cost-Benefit Trade-off:** The correction loop increases API calls by an average of 15%, but provides meaningful accuracy gains.

## 6 Debug Diary

During reproduction, I encountered and resolved several technical challenges:

### 6.1 Issue 1: Missing Module

**Problem:** The repository lacked `analyze_by_subproblems.py`, causing import errors.

**Root Cause:** Module not committed to GitHub repository.

**Solution:** Implemented the missing module by reverse-engineering its expected interface from usage patterns in the main script. My implementation supports multiple JSON formats from the Subproblem Agent.

### 6.2 Issue 2: Spider Dataset Path

**Problem:** Hardcoded relative paths (`../../spider/`) did not match my directory structure.

**Solution:** Updated all Spider references to `../spider/` and verified database file accessibility.

### 6.3 Issue 3: SQLite Version Compatibility

**Problem:** Some queries used SQLite-specific functions not available in older versions.

**Verification:** Confirmed Spider databases use SQLite 3.36+. Updated system SQLite from 3.31 to 3.40.

## 7 Conclusions

### 7.1 Reproducibility Assessment

Table 5: Reproducibility Summary

Aspect	Status	Notes
Code Availability	Full	GitHub repository accessible
Environment Setup	Straightforward	Minimal dependencies
Data Availability	Public	Spider dataset freely available
Documentation	Partial	Missing module implementations
Results Reproducible	Yes	Within expected model variance

### 7.2 Key Findings

- Architecture Validated:** The multi-agent approach successfully improves Text-to-SQL performance over single-model baselines, even with a weaker LLM.
- Model Dependency:** Results are sensitive to LLM capability; the 10.74% performance gap between DeepSeek and GPT-4o-mini highlights the importance of model selection.
- Correction Loop Value:** The error correction mechanism provides meaningful improvements (6.2%) at modest computational cost.
- Practical Viability:** The system runs efficiently (15s per sample) and could be deployed in production with appropriate model selection.

### 7.3 Limitations

1. **Model Substitution:** Direct comparison with paper results is limited due to LLM differences.
2. **Compute Budget:** Using other more powerful LLMs (like Gemini and Claude) may perform better but would require significant API costs.
3. **Stochasticity:** With temperature=0.0, results are deterministic, but variance analysis would require multiple runs with sampling.

### 7.4 Recommendations for Future Users

1. **Model Selection:** Use GPT-5.2 or Claude Opus for closer reproduction of original results
2. **Error Handling:** Implement robust JSON parsing with fallback mechanisms for malformed agent outputs
3. **Logging:** Add detailed logging of agent outputs for debugging and error analysis

### 7.5 Contributions to Community

The reproduction study contributes:

- Validation of multi-agent architecture effectiveness
- Quantification of correction loop contribution
- Identification of common failure modes in smaller LLMs

## References

- [1] [Original Authors]. *SQL-of-Thought: Multi-agentic Text-to-SQL with Guided Error Correction*. NeurIPS 2025 Workshop on Deep Learning for Code. <https://arxiv.org/abs/2509.00581>
- [2] Yu, T., Zhang, R., Yang, K., Yasunaga, M., Wang, D., Li, Z., ... & Radev, D. (2018). *Spider: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-SQL task*. EMNLP 2018.
- [3] DeepSeek. (2024). *DeepSeek API Documentation*. <https://platform.deepseek.com/api-docs/>

## A Code Repository

Complete code and results available at: <https://github.com/Steven-Zhan/FTEC5660/tree/main/individual%20project/SQL-of-Thought-main>

Repository includes:

- Modified source code with clear commit history
- README with setup and reproduction instructions
- Results JSON files for all experiments
- Requirements file with exact package versions

## B Sample Output

Listing 1: Example Unsuccessful Generation

```
1 Question: Show the property type descriptions of properties belonging
2   to that code.
3 Database: real_estate_properties
4
5 [1/5] Schema Linking Agent...
6 [Completed] Output Length: 213 characters
7
8 [2/5] Subproblem Agent...
9 [Completed] Output: {
10   "clauses": ["SELECT", "JOIN"],
11   "subproblems": [
12     {"clause": "SELECT", "description": ...
13   [Identified SQL Clauses]: ['JOIN', 'SELECT']
14
15 [3/5] Query Plan Agent...
16 [Completed] Plan: **Query Plan:**

17 1. **Tables to use:**
18    - 'Ref_Property_Types' (contains property type descriptions)
19    - 'Properties' (contains property record...
20
21 [4/5] SQL Agent...
22 [Generated SQL]: select ref_property_types.property_type_description
23   from properties inner join ref_property_types on properties.
24   property_type_code = ref_property_types.property_type_code
25
26 [5/5] SQL Correction Loop (Error: None)
27
28 Correction Attempt 1/3
29 [Correction Plan]: **Correction Plan:**

30 The SQL is syntactically correct but does not match the questions
31 requirement.
32 [Corrected SQL]: select ref_property_types.property_type_description
33   from properties inner join ref_property_types on properties.
   property_type_code = ref_property_types.property_type_code where
   properties.property_type_code = ?
```

```

34
35 1. **Identify the Issue:** The SQL query uses a parameter placeholder
   ('?')
36 [Corrected SQL]: select ref_property_types.property_type_description
   from properties inner join ref_property_types on properties.
   property_type_code = ref_property_types.property_type_code
37
38 Correction Attempt 3/3
39 [Correction Plan]: **Correction Plan:***
40
41 The SQL is syntactically correct and should run without error.
42 [Corrected SQL]: select ref_property_types.property_type_description
   from properties inner join ref_property_types on properties.
   property_type_code = ref_property_types.property_type_code where
   properties.property_type_code = 'x'
43   Maximum number of attempts reached
44
45 Exact Match: False
46 Valid SQL: True
47 Execution Accuracy: False

```

Listing 2: Example Successful Generation

```

1 Question: What are the names of properties that are either houses or
   apartments with more than 1 room?
2 Database: real_estate_properties
3
4 [1/5] Schema Linking Agent...
5 [Completed] Output Length: 242 characters
6
7 [2/5] Subproblem Agent...
8 [Completed] Output: {
9   "clauses": ["SELECT", "WHERE", "JOIN"],
10  "subproblems": [
11    {"clause": "SELECT", "des...
12 [Identified SQL Clauses]: ['JOIN', 'WHERE', 'SELECT']
13
14 [3/5] Query Plan Agent...
15 [Completed] Plan: **Query Plan:***
16
17 1. **Tables to use:***
18   - 'Properties' (main table for property details)
19   - 'Ref_Property_Types' (to identify property types by default)
20
21 [4/5] SQL Agent...
22 [Generated SQL]: select p.property_name from properties p inner join
   ref_property_types r on p.property_type_code = r.property_type_code
   where r.property_type_description in ('house', 'apartment') and p.
   room_count > 1
23
24 Exact Match: False
25 Valid SQL: True
26 Execution Accuracy: True

```